RECURRING POLICY DECISIONS

# Canopy makes a verdict inspectable

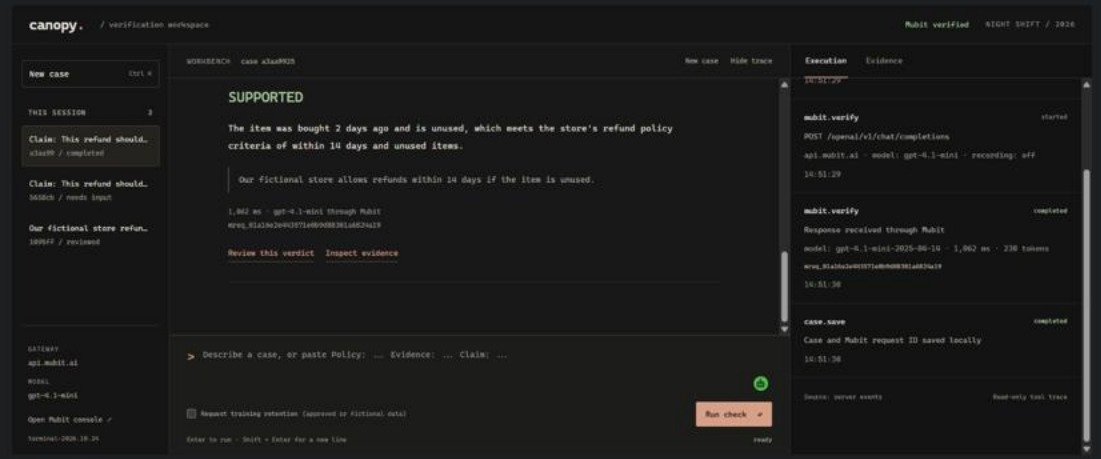
A reviewer can see the rule, the supplied facts, and the Mubit trace.

THE JOB

## Policy + evidence + claim

A clear verdict, its reason, and
an exact supplied policy quote.

SUPPORTED  VIOLATES
INSUFFICIENT




Earlier live fictional-case capture

LIVE EXAMPLE

# The damaged-item exception wins

The built-in refund case is fictional and safe to show in public.

## SUPPLIED POLICY

Damaged items: refund within 30 days,
even after use. Normal window: 14 days.
Evidence: 21 days, used, arrived damaged,
photos supplied.

## CLAIM AND VERDICT

Claim: refuse because used and late.
Canopy verdict: VIOLATES.
Reason: the damaged-item exception
covers the supplied facts.

A reviewer confirmed this correct verdict; the review is saved locally.

MUBIT INTEGRATION

# One traceable inference path

Extraction, verification, and chat use the Mubit OpenAI gateway.




Live verification response ID:

`mreq_01a16ee996087620822b6764c56cf1dd`

A reviewer confirmation remained visible after browser reload.

JUDGING EVIDENCE

# Learning status without invention

Mubit support accepts app routing plus a labelled Try models comparison.

## PROVEN NOW

Live Mubit-routed verification
Request IDs and policy evidence
Review persists after browser reload
35 local tests; hosted CI passed

## OPEN LEARNING GATE

Review remains local_only
Kept examples need console proof
Try models calls do not train
Compare only a Passed adapter

If no adapter passes in time, show the current task status and say so.

Code: github.com/ShieldTech-Ltd/canopy